

Notebook #1 Image classification with CNN

Input: 32x32 RGB images in 10 classes from the CIFAR-10 dataset

Output: Predicted class probabilities or labels for the 10 categories.

Hardware requirement: NVIDIA GPUs

Process

1. Preprocess

- 1.1. Install and import libraries
- 1.2. Load datasets from CIFAR-10
- 1.3. Apply normalization and data augmentation

2. Create CNN model

2.1. Pipeline

- First convolutional layer has 3 input and 6 output channels with 5x5 kernel.
- Pooling layer uses a 2x2 kernel with a stride of 2.
- Second convolutional layer has 6 input and 16 output channels with 5x5 kernel.
- First fully connected layer takes 400 input and outputs 120 units.
- Second fully connected layer takes 120 input units and outputs 84 units.
- Third fully connected layer takes 84 input units and outputs 10 units

2.2. Summary

- Activation function: ReLU, Softmax
- Loss function: Cross-Entropy Loss
- Optimizer: Adam
- Total params: 62,006

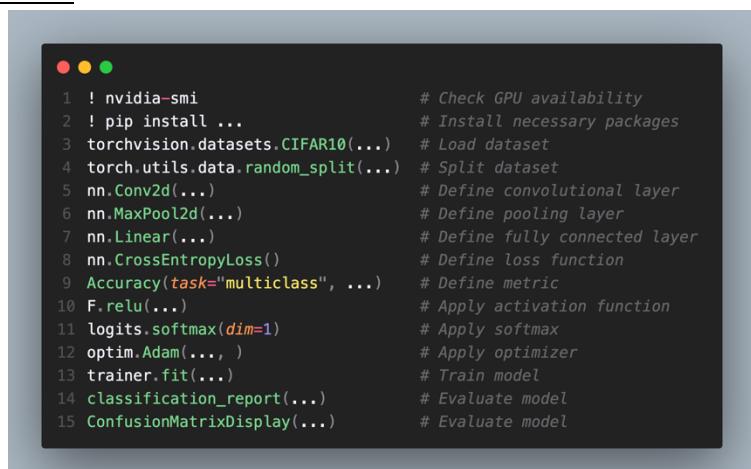
3. Training

- Batch size: 32
- Epoch: 1
- 60,000 images split into: 50,000 training images and 10,000 test images
- Each class contains an equal number of images.

4. Evaluate the model

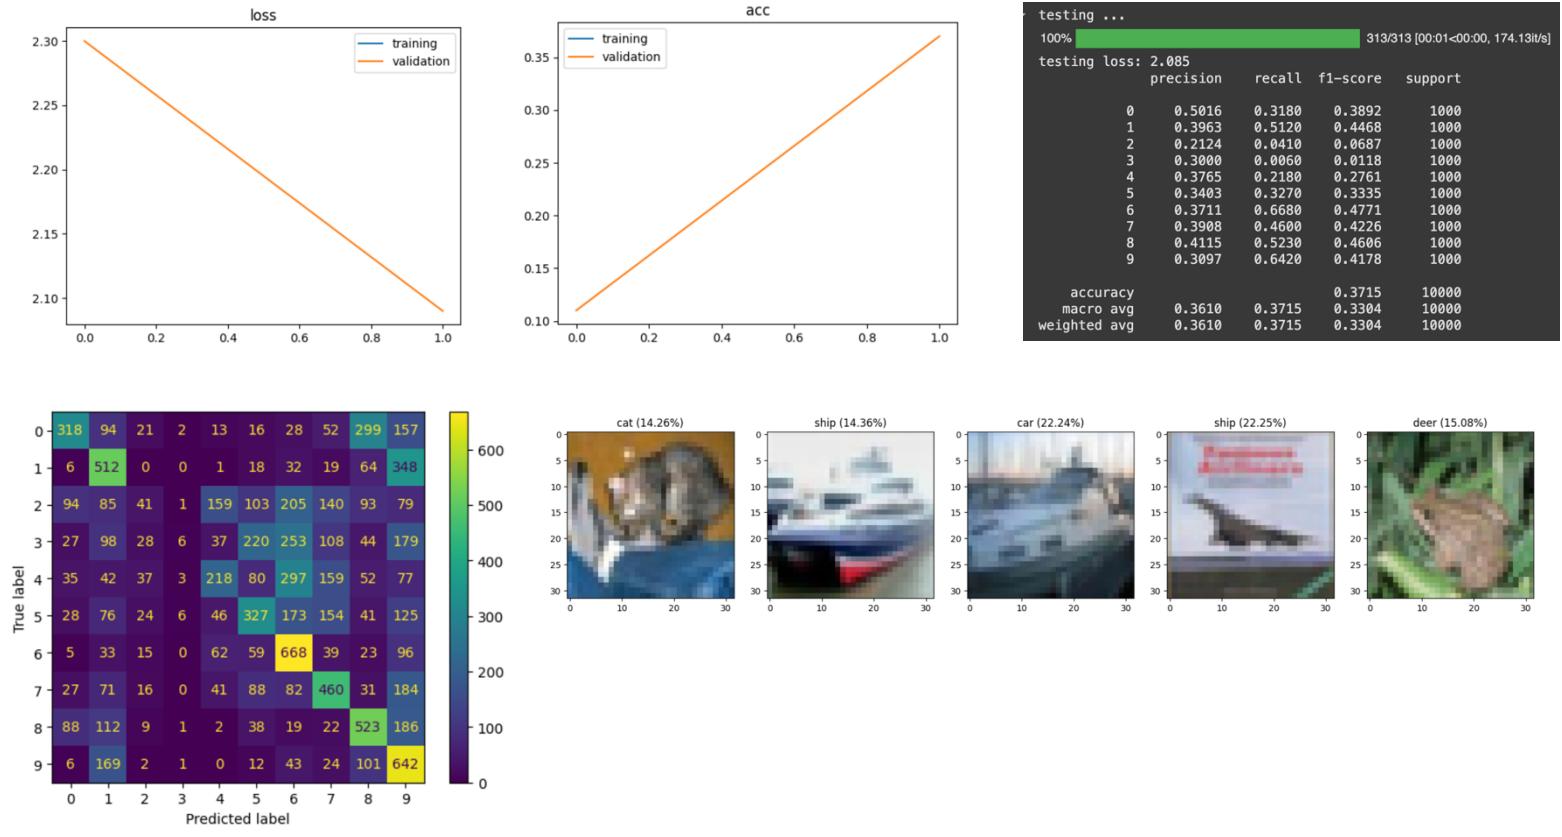
- Accuracy, Confusion Matrix

Important commands



```
! nvidia-smi # Check GPU availability
! pip install ... # Install necessary packages
torchvision.datasets.CIFAR10(...) # Load dataset
torch.utils.data.random_split(...) # Split dataset
nn.Conv2d(...) # Define convolutional layer
nn.MaxPool2d(...) # Define pooling layer
nn.Linear(...) # Define fully connected layer
nn.CrossEntropyLoss() # Define loss function
Accuracy(task="multiclass", ...) # Define metric
F.relu(...) # Apply activation function
logits.softmax(dim=1) # Apply softmax
optim.Adam(..., ) # Apply optimizer
trainer.fit(...) # Train model
classification_report(...) # Evaluate model
ConfusionMatrixDisplay(...) # Evaluate model
```

Results



Notebook #2 Image classification with EfficientNet

Input: Image dataset with 10 animal classes

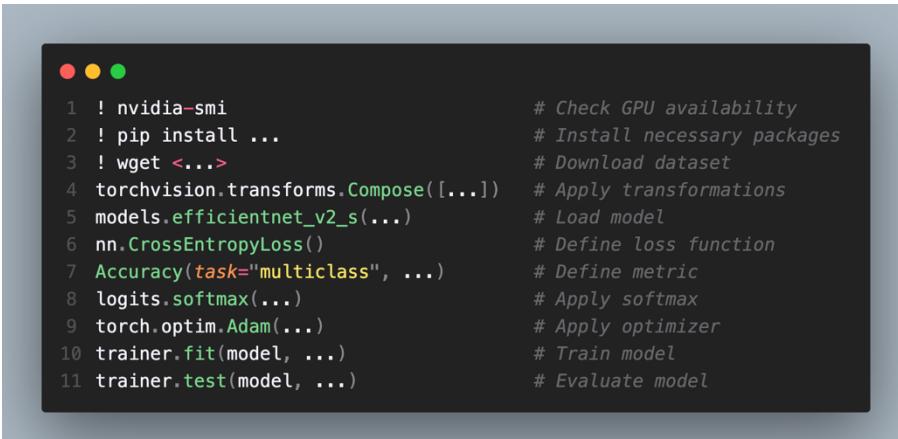
Output: Predicted class labels

Hardware requirement: NVIDIA GPUs

Process

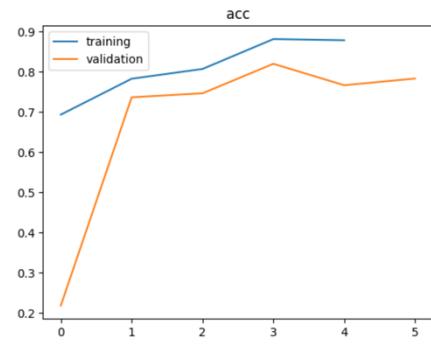
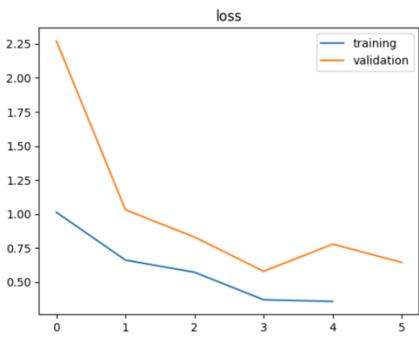
1. Preprocess
 - 1.1. Install and import libraries
 - 1.2. Load datasets from GitHub repository
 - 1.3. Image augmentation ex. Rotation, Horizontal and Vertical flip
 - 1.4. Normalize images
 - 1.5. Split data into 1,400 training images, 300 validation images and 300 test images
2. Training Setup
 - Model: EfficientNetV2
 - Loss function: Cross-Entropy function
 - Optimizer: Adam
 - Learning rate: 0.001
 - Batch size: 32
 - Epoch: 5
 - Trainable parameters: 20.2 M
3. Training
4. Evaluate the model
 - Accuracy, Confusion Matrix

Important Commands

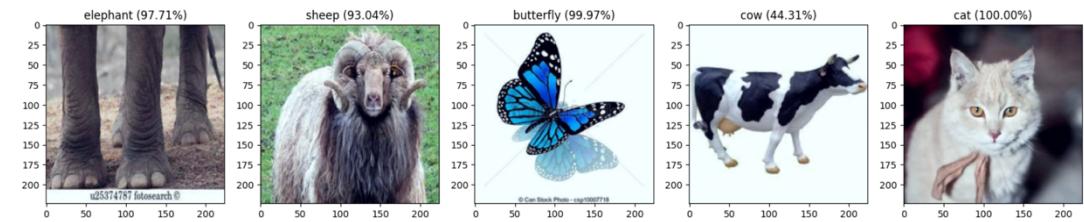
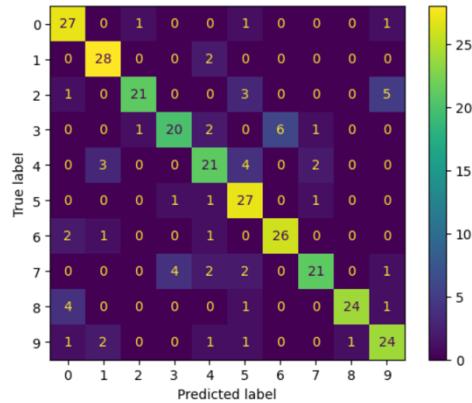


```
! nvidia-smi # Check GPU availability
! pip install ... # Install necessary packages
! wget <....> # Download dataset
torchvision.transforms.Compose([...]) # Apply transformations
models.efficientnet_v2_s(...) # Load model
nn.CrossEntropyLoss() # Define loss function
Accuracy(task="multiclass", ...) # Define metric
logits.softmax(...) # Apply softmax
torch.optim.Adam(...) # Apply optimizer
trainer.fit(model, ...) # Train model
trainer.test(model, ...) # Evaluate model
```

Results



testing ...				
100% [00:53<00:00, 4.30s/it]				
testing loss: 0.7089				
precision	recall	f1-score	support	
0 0.7714	0.9000	0.8308	30	
1 0.8235	0.9333	0.8750	30	
2 0.9130	0.7000	0.7925	30	
3 0.8000	0.6667	0.7273	30	
4 0.7000	0.7000	0.7000	30	
5 0.6923	0.9000	0.7826	30	
6 0.8125	0.8667	0.8387	30	
7 0.8400	0.7000	0.7636	30	
8 0.9600	0.8000	0.8727	30	
9 0.7500	0.8000	0.7742	30	
accuracy			0.7967	300
macro avg	0.8063	0.7967	0.7957	300
weighted avg	0.8063	0.7967	0.7957	300



Notebook #3 Object detection with yolov8 (basic)

Input: Images in 20 classes from the Pascal-VOC dataset.

Output: Predicted class labels

Hardware requirement: NVIDIA GPUs

Process:

1. Preprocess
 - 1.1. Import packages from Ultralytics
 - 1.2. Checks software and hardware configuration
2. Train
 - Model: YOLOv8
 - Images of a size 640x640 with 20 classes
 - Epoch: 3
 - Learning rate: 0.0004
 - Optimizer: Adam
 - Parameters: 3 M
3. Evaluate model

Important commands

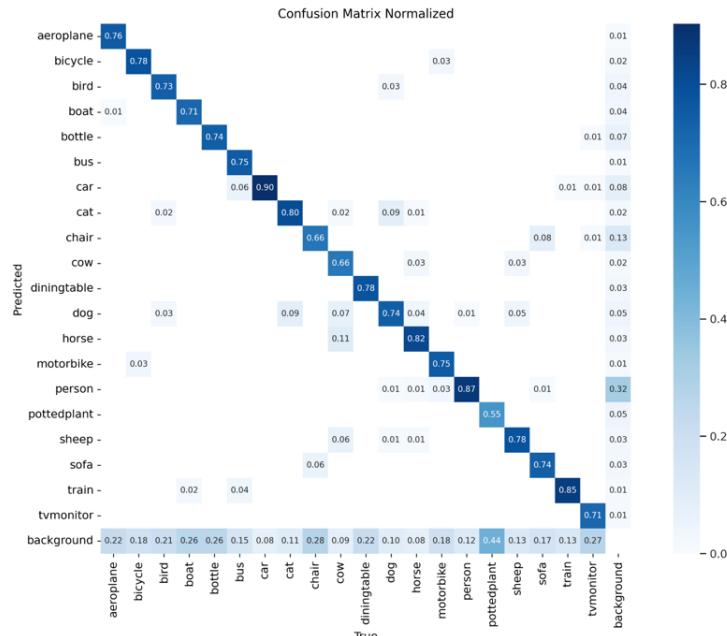
```
 1 %pip install ultralytics                                         # Install Ultralytics
 2 !nvidia-smi                                                 # Check GPU availability
 3
 4 # For CLI usage
 5 !yolo train data=VOC.yaml model=yolov8n.pt epochs=3 imgsze=640 device=0  # Train model
 6 !yolo predict model=yolov8n.pt source='image.jpg' device=0           # Predict image
 7
 8 # For Python usage
 9 model = YOLO(...)

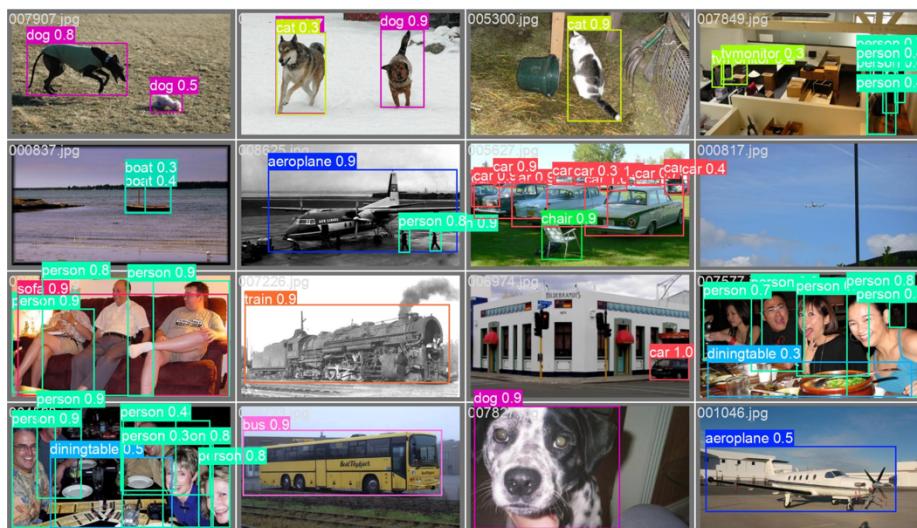
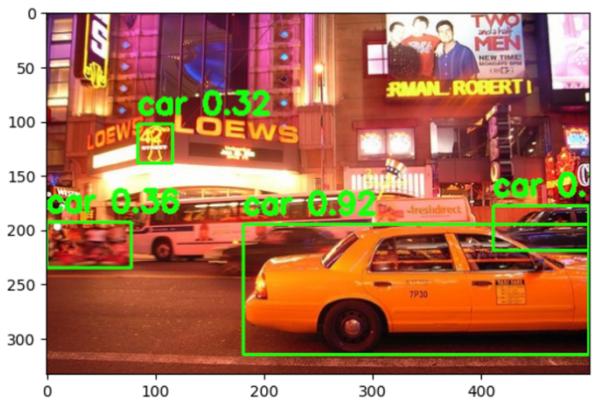
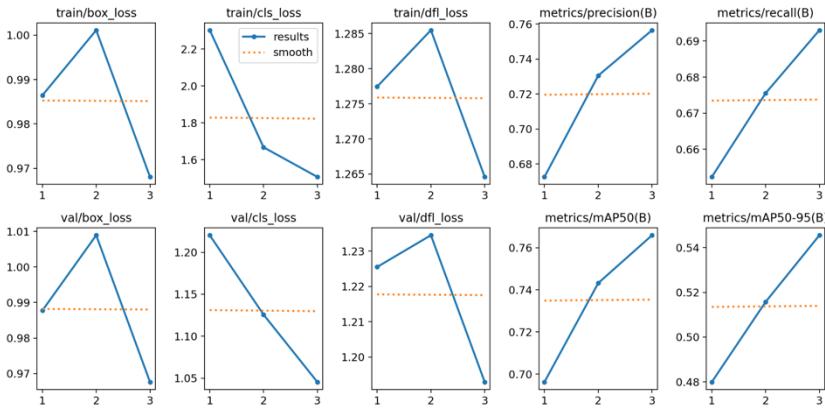
10 model.train(...)

11 model.val()

12 model.predict(...)
```

Results





Notebook #4 Object detection with yolov8 (advanced)

Input: Pascal VOC datasets of 20 object classes

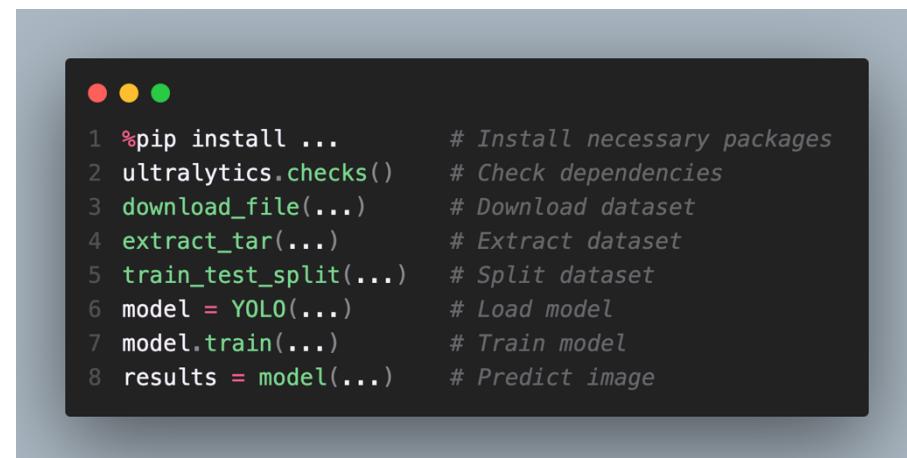
Output: Trained YOLOv8 object detection model with bounding box predictions

Hardware requirement: NVIDIA GPUs

Process:

1. Preprocess
 - 1.1. Install and import packages
 - 1.2. Download and extract Pascal VOC dataset
 - 1.3. Define .yml file structure for YOLO
 - 1.4. Convert VOC XML annotations to YOLO format
2. Setup model
 - Model: YOLOv8
 - Epoch: 3
 - Learning rate: 0.0004
 - Optimizer: Adam
 - Parameters: 3 M
 - 10,000 images split into 8,000 training images and 2,000 test images
3. Train
 - 3.1. Load YOLOv8 model
 - 3.2. Train with dataset and monitor loss
4. Evaluate
 - 4.1. Validate performance with mAP
5. Test
 - 5.1. Test detection on sample images

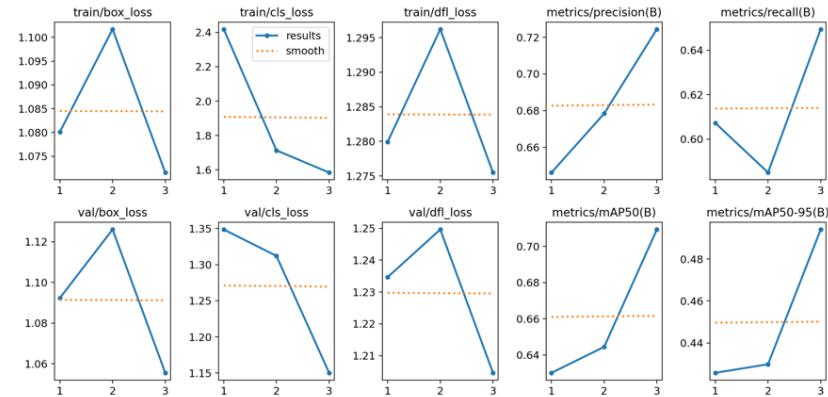
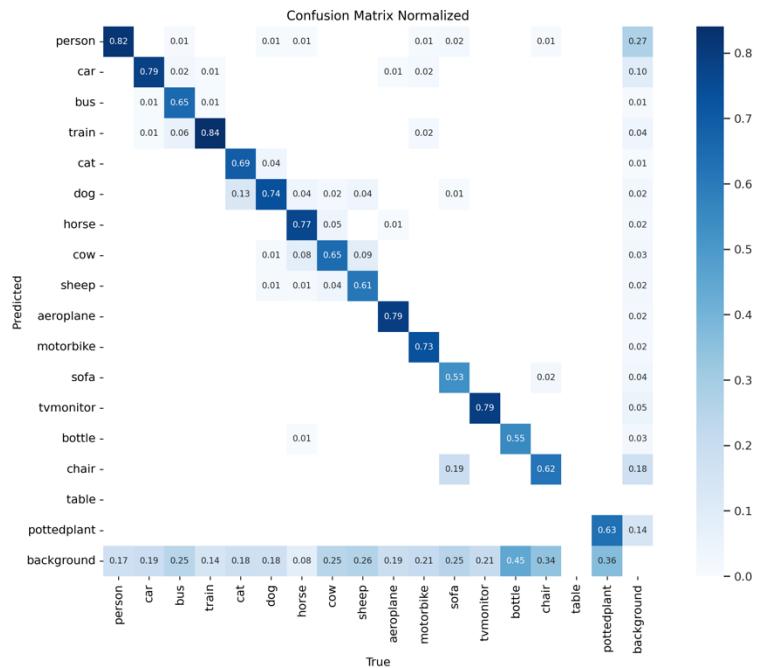
Important commands



```
● ● ●

1 %pip install ...          # Install necessary packages
2 ultralytics.checks()       # Check dependencies
3 download_file(...)         # Download dataset
4 extract_tar(...)           # Extract dataset
5 train_test_split(...)      # Split dataset
6 model = YOLO(...)          # Load model
7 model.train(...)           # Train model
8 results = model(...)        # Predict image
```

Results



Notebook #5 Semantic segmentation with DeepLabV3

Input: Image from CamSeq dataset with a size of 960x720

Output: Segmentation mask

Hardware requirement: NVIDIA GPUs

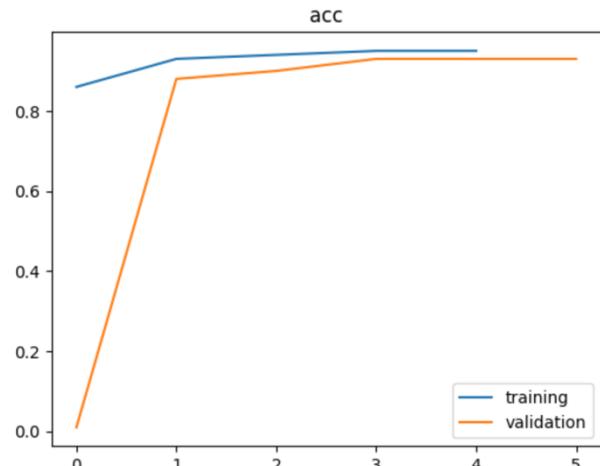
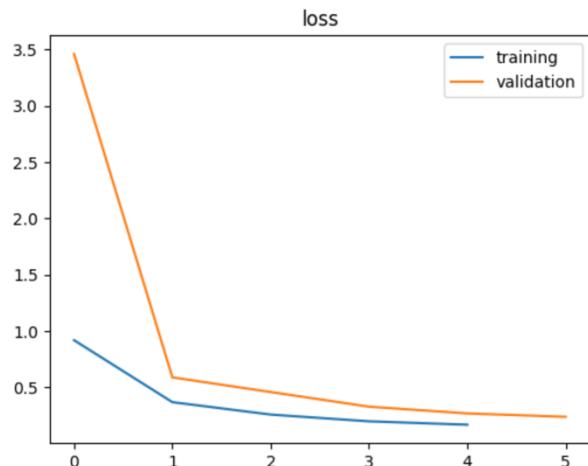
Process:

1. Preprocess
 - 1.1. Install and import packages
 - 1.2. Download 100 image datasets from CamSeq
 - 1.3. Load colormap from label_color.txt
 - 1.4. Create Data loader
 - 1.5. Apply transformations (resize, normalize).
 - 1.6. Split 101 images into 70 training images, 15 validation images and 16 test images
2. Setup model
 - Model: DeepLabV3
 - Trainable parameters: 42 M
 - Epoch: 2
 - Training Metrics: Cross-Entropy, accuracy and IoU.
3. Train model
4. Evaluate model
 - Compute IoU
5. Visualize results
 - Test model with sample image

Important commands

```
● ● ●
1 ! wget <....>                                # Download dataset
2 ! pip install ...                               # Install necessary packages
3 ! nvidia-smi                                 # Check GPU availability
4 ! unzip <...>                                 # Extract dataset
5 A.Normalize(...)                                # Apply transformations
6 trainer = pl.Trainer(...)                      # Initialize trainer
7 trainer.fit(...)                                # Train model
8 model = deeplabv3.load_from_checkpoint(...)    # Load model
9 model.eval()                                    # Evaluate model
10 imshow(image)                                 # Display image
```

Results:



```
100% [8/8 00:16<00:00, 2.08s/it]
average_iou_per_class: tensor([ nan, 0.000, 0.6595,    nan, 0.7270, 0.0045,    nan, 0.3903, 0.0297,
    0.7489, 0.0631,    nan, 0.000,    nan, 0.1779,    nan, 0.0212, 0.8938,
    nan, 0.3790, 0.000, 0.7314,    nan,    nan, 0.4691,    nan, 0.5835,
    0.4388,    nan, 0.0012, 0.0679, 0.5700])
overall_average_iou: 0.33127713203430176
```

